

Learning Kernels with Random Features

Aman Sinha & John Duchi, Stanford University
{amans, jduchi}@stanford.edu

Kaikai Zhao
2017/10/8

Introduction

- An essential element of supervised learning systems is **the representation of input data**.
- Kernel methods provide one approach to this problem: they **implicitly** transform the data to a new feature space, **allowing non-linear data representations**.
- This representation comes with a **cost**.

$$K(x_i, x_j) = \exp(-\|x_i - x_j\|^2 / 2\sigma^2), \quad \sigma \text{ width}$$

Introduction

- Instead of implicit representation, [1] proposes explicitly mapping the data to a low-dimensional Euclidean inner product space using a randomized feature map \mathbf{z} :
- $\mathbb{R}^d \rightarrow \mathbb{R}^D$ approximates their kernel evaluation:

$$k(\mathbf{x}, \mathbf{y}) = \langle \phi(\mathbf{x}), \phi(\mathbf{y}) \rangle \approx \mathbf{z}(\mathbf{x})' \mathbf{z}(\mathbf{y}).$$

Merits ???

1. Ali Rahimi, Benjamin Recht. Random Features for Large-Scale Kernel Machines. NIPS, 2007

Introduction

- With the kernel trick, evaluating the machine at a test point \mathbf{x} requires computing

$$f(\mathbf{x}) = \sum_{i=1}^N c_i k(\mathbf{x}_i, \mathbf{x})$$

- which requires $O(Nd)$ operations to compute and requires retaining much of the dataset unless the machine is very sparse.
- On the other hand, after learning a hyperplane \mathbf{w} , a linear machine can be evaluated by simply computing

$$f(x) = \mathbf{w}'\mathbf{z}(\mathbf{x})$$




- which, with the randomized feature maps presented here, requires only $O(D + d)$ operations and storage.

傅里叶变换公式

$$F(\omega) = \mathcal{F}[f(t)] = \int_{-\infty}^{\infty} f(t)e^{-i\omega t} dt$$

公式描述： 公式中 $F(\omega)$ 为 $f(t)$ 的像函数， $f(t)$ 为 $F(\omega)$ 的像原函数。

$$f(t) = \mathcal{F}^{-1}[F(\omega)] = \frac{1}{2\pi} \int_{-\infty}^{\infty} F(\omega)e^{i\omega t} d\omega$$



Random Fourier Features

Theorem 1 (Bochner [15]). *A continuous kernel $k(x, y) = k(x - y)$ on \mathcal{R}^d is positive definite if and only if $k(\delta)$ is the Fourier transform of a non-negative measure.*




Figure 1: Random Fourier Features. Each component of the feature map $\mathbf{z}(\mathbf{x})$ projects \mathbf{x} onto a random direction ω drawn from the Fourier transform $p(\omega)$ of $k(\Delta)$, and wraps this line onto the unit circle in \mathcal{R}^2 . After transforming two points \mathbf{x} and \mathbf{y} in this way, their inner product is an unbiased estimator of $k(\mathbf{x}, \mathbf{y})$. The table lists some popular shift-invariant kernels and their Fourier transforms. To deal with non-isotropic kernels, the data may be whitened before applying one of these kernels.

Random Fourier Features

If the kernel $k(\delta)$ is properly scaled, Bochner's theorem guarantees that its Fourier transform $p(\omega)$ is a proper probability distribution. Defining $\zeta_\omega(\mathbf{x}) = e^{j\omega' \mathbf{x}}$, we have

$$k(\mathbf{x} - \mathbf{y}) = \int_{\mathcal{R}^d} p(\omega) e^{j\omega'(\mathbf{x} - \mathbf{y})} d\omega = E_\omega[\zeta_\omega(\mathbf{x})\zeta_\omega(\mathbf{y})^*], \quad (2)$$

so $\zeta_\omega(\mathbf{x})\zeta_\omega(\mathbf{y})^*$ is an unbiased estimate of $k(\mathbf{x}, \mathbf{y})$ when ω is drawn from p .

To obtain a real-valued random feature for k , note that both the probability distribution $p(\omega)$ and the kernel $k(\Delta)$ are real, so the integrand $e^{j\omega'(\mathbf{x} - \mathbf{y})}$ may be replaced with $\cos \omega'(\mathbf{x} - \mathbf{y})$. Defining $z_\omega(\mathbf{x}) = [\cos(\omega' \mathbf{x}) \sin(\omega' \mathbf{x})]'$ gives a real-valued mapping that satisfies the condition $E[z_\omega(\mathbf{x})' z_\omega(\mathbf{y})] = k(\mathbf{x}, \mathbf{y})$, since $z_\omega(\mathbf{x})' z_\omega(\mathbf{y}) = \cos \omega'(\mathbf{x} - \mathbf{y})$. Other mappings such as $z_\omega(\mathbf{x}) = \sqrt{2} \cos(\omega' \mathbf{x} + b)$, where ω is drawn from $p(\omega)$ and b is drawn uniformly from $[0, 2\pi]$, also satisfy the condition $E[z_\omega(\mathbf{x})' z_\omega(\mathbf{y})] = k(\mathbf{x}, \mathbf{y})$.

$$e^{(-j\omega t)} = \cos(\omega t) - j\sin(\omega t)$$

$$e^{(j\omega t)} = \cos(\omega t) + j\sin(\omega t)$$

Problem Setup and Approach

Learning a kernel

Input: n datapoints $(x^i, y^i) \in \mathbb{R}^d \times \{-1, 1\}$

Problem: Consider a kernel $K_Q(x, x') = \mathbb{E}_Q[\phi(x; W)\phi(x'; W)]$. We want a good Q

- Try kernel alignment: maximize $_{Q \in \mathcal{P}} \sum_{i,j} K_Q(x^i, x^j) y^i y^j$
- Constrain Q with power f-divergences around a base distribution:

$$f(t) = t^k - 1 \ (k \geq 2), \quad \mathcal{P} := \{Q : D_f(Q \| P_0) \leq \rho\}$$

- Approximate the problem with random features:

$$\underset{q \in \mathcal{P}_{N_w}}{\text{maximize}} \quad \sum_{i,j} y^i y^j \sum_{m=1}^{N_w} q_m \phi(x^i, w^m) \phi(x^j, w^m) \quad (1)$$

where $w^i \stackrel{\text{iid}}{\sim} P_0$ and $\mathcal{P}_{N_w} := \{q : D_f(q \| \mathbf{1}/N_w) \leq \rho\}$

- Fast (near-linear time) solution, often results in sparse q
- At a high level, we take a feature mapping, find a distribution that aligns this mapping with the labels \mathbf{y} , and draw random features from the learned distribution; we then use these features in a standard supervised learning approach.

Efficiently solving problem

$$\Phi = [\phi^1 \ \dots \ \phi^{\overline{n}}] \in \mathbb{R}^{N_w \times n} \quad \phi^i = [\phi(x^i, w^1) \ \dots \ \phi(x^i, w^{N_w})]^T \in \mathbb{R}^{N_w}$$

- is the randomized feature representation for x^i and $w^m \sim P_0$, we can rewrite the optimization objective as

$$\sum_{i,j} y^i y^j \sum_{m=1}^{N_w} q_m \phi(x^i, w^m) \phi(x^j, w^m) = \sum_{m=1}^{N_w} q_m \left(\sum_{i=1}^n y^i \phi(x^i, w^m) \right)^2 = q^T ((\Phi y) \odot (\Phi y))$$

$$\overline{D}_f(Q \| P_0) \leq \rho$$

Lagrangian L:

$$\mathcal{L}(q, \lambda) = q^T ((\Phi y) \odot (\Phi y)) - \lambda (D_f(q \| \mathbf{1}/N_w) - \rho)$$

Lagrange dual function:

$$g(\lambda) = \sup_{q \in \Delta} \mathcal{L}(q, \lambda)$$

Probability simplex:

$$\Delta := \{q \in \mathbb{R}_{+}^{\overline{N_w}} : q^T \mathbf{1} = 1\}$$

Efficiently solving problem

- Minimizing $g(\lambda)$ yields the solution to problem
- This is a **convex optimization problem in one dimension** so we can use **bisection**
- Each iteration is maximizing $L(q, \lambda)$ with respect to q for **a given λ**

- For $f(t) = t^k - 1$, we define $v := (\Phi y) \odot (\Phi y)$

- solve
$$\underset{q \in \Delta}{\text{maximize}} \quad q^T v - \lambda \frac{1}{N_w} \sum_{m=1}^{N_w} (N_w q_m)^k.$$

- KKT conditions

$$q_m = \left[v_m / \lambda N_w^{k-1} + \tau \right]_+^{\frac{1}{k-1}} \quad \sum_m q_m = 1.$$

$$\underset{q \in \Delta}{\text{maximize}} \quad q^T v - \lambda \frac{1}{N_w} \sum_{m=1}^{N_w} (N_w q_m)^k. \quad (8)$$

$$q_m = \left[v_m / \lambda N_w^{k-1} + \tau \right]_+^{\frac{1}{k-1}} \quad \sum_m q_m = 1.$$

A small typo in the form of q_m
A **missing k** in the denominator

Algorithm 1 Kernel optimization with $f(t) = t^k - 1$ as divergence

INPUT: distribution P_0 on \mathcal{W} , sample $\{(x^i, y^i)\}_{i=1}^n$, $N_w \in \mathbb{N}$, feature function ϕ , $\epsilon > 0$

OUTPUT: $q \in \mathbb{R}^{N_w}$ that is an ϵ -suboptimal solution to (4).

SETUP: Draw N_w samples $w^m \stackrel{\text{iid}}{\sim} P_0$, build feature matrix Φ , compute $v := (\Phi y) \odot (\Phi y)$.

Set $\lambda_u \leftarrow \infty$, $\lambda_l \leftarrow 0$, $\lambda_s \leftarrow 1$

while $\lambda_u = \infty$

$q \leftarrow \arg\max_{q \in \Delta} \mathcal{L}(q, \lambda_s)$ // (solution to problem (8))

if $D_f(q \| \mathbf{1}/N_w) < \rho$ **then** $\lambda_u \leftarrow \lambda_s$ **else** $\lambda_s \leftarrow 2\lambda_s$

while $\lambda_u - \lambda_l > \epsilon \lambda_s$

$\lambda \leftarrow (\lambda_u + \lambda_l)/2$

$q \leftarrow \arg\max_{q \in \Delta} \mathcal{L}(q, \lambda)$ // (solution to problem (8))

if $D_f(q \| \mathbf{1}/N_w) < \rho$ **then** $\lambda_u \leftarrow \lambda$ **else** $\lambda_l \leftarrow \lambda$

However, it is not the case in the MATLAB implementation, which is confusing.

$$\chi^2(P|Q) = \sum_{i=1}^n \frac{(p_i - q_i)^2}{q_i} \approx \sum_{i=1}^n \frac{p_i^2}{q_i} - 1$$

project_onto_simplex

$$\underset{\mathbf{w}}{\text{minimize}} \frac{1}{2} \|\mathbf{w} - \mathbf{v}\|_2^2 \quad \text{s.t.} \quad \sum_{i=1}^n w_i = z, \quad w_i \geq 0$$

When $z = 1$ the above is projection onto the probabilistic simplex. The Lagrangian of the problem in Eq. (3) is

$$\mathcal{L}(\mathbf{w}, \zeta) = \frac{1}{2} \|\mathbf{w} - \mathbf{v}\|^2 + \theta \left(\sum_{i=1}^n w_i - z \right) - \zeta \cdot \mathbf{w},$$

where $\theta \in \mathbb{R}$ is a Lagrange multiplier and $\zeta \in \mathbb{R}_+^n$ is a vector of non-negative Lagrange multipliers. Differentiating with respect to w_i and comparing to zero gives the optimality condition, $\frac{d\mathcal{L}}{dw_i} = w_i - v_i + \theta - \zeta_i = 0$. The complementary slackness KKT condition implies that whenever $w_i > 0$ we must have that $\zeta_i = 0$. Thus, if $w_i > 0$ we get that

$$w_i = v_i - \theta + \zeta_i = v_i - \theta. \quad (4)$$

INPUT: A vector $\mathbf{v} \in \mathbb{R}^n$ and a scalar $z > 0$

Sort \mathbf{v} into $\boldsymbol{\mu} : \mu_1 \geq \mu_2 \geq \dots \geq \mu_p$

Find $\rho = \max \left\{ j \in [n] : \mu_j - \frac{1}{j} \left(\sum_{r=1}^j \mu_r - z \right) > 0 \right\}$

Define $\theta = \frac{1}{\rho} \left(\sum_{i=1}^{\rho} \mu_i - z \right)$

OUTPUT: \mathbf{w} s.t. $w_i = \max \{ v_i - \theta, 0 \}$

Figure 1. Algorithm for projection onto the simplex.

% min. $x' * v$

- % s.t. $\text{norm}(x - u, 2)^2 \leq \rho$

- % $\text{sum}(x) == 1, x \geq 0.$

- % A partial dual to the problem is given by the Lagrangian

- %

- % $L(x, \lambda) = (\lambda/2) * (\text{norm}(x - u, 2)^2 - \rho) + x' * v$

- % subject to $\text{sum}(x) == 1, x \geq 0.$

- %

Consistency and generalization performance guarantees

$$\underset{q \in \mathcal{P}_{N_w}}{\text{maximize}} \quad \sum_{i,j} y^i y^j \sum_{m=1}^{N_w} q_m \phi(x^i, w^m) \phi(x^j, w^m). \quad (4)$$

- Although the procedure (4) is a discrete approximation to a heuristic kernel alignment problem, we can provide guarantees on its **consistency** as well as the **generalization** performance of our subsequent model trained with the optimized kernel.
- **Three lemmas and two theorems (proofs in the supplement)**

Empirical evaluations

- synthetic data $x^i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, I)$
- Labels $y^i = \text{sign}(\|x\|_2 - \sqrt{d}), \quad x \in \mathbb{R}^d.$
- Gaussian kernel corresponds to $\bar{\phi}(x, (w, v)) = \cos((x, 1)^T(w, v))$
- Training set $n=10^4$, test set 10^3 $(W, V) \sim \mathcal{N}(0, I) \times \text{Uni}(0, 2\pi)$
- The Gaussian kernel is **ill-suited** for this task, as the **Euclidean distance** used in this kernel does not capture the underlying structure of the classes.
- Random features $N_w = 2 \cdot 10^4$





Figure 1 shows the results of the experiments for $d \in \{2, \dots, 15\}$. Figure 1(a) illustrates the output of the optimization when $d = 2$. The selected kernel features w^m lie near $(1, 1)$ and $(-1, -1)$; the offsets v^m are near 0 and π , giving the feature $\phi(\cdot, w, v)$ a parity flip. Thus, the kernel computes similarity between datapoints via neighborhoods of $(1, 1)$ and $(-1, -1)$ close to the classification boundary. In higher dimensions, this generalizes to neighborhoods of pairs of opposing points along the surface of the d -sphere; these features provide a coarse approximation to vector magnitude.

